

RAKSHITH S

9741913621 | rakshiths2001official@gmail.com | [LinkedIn](#) | [Github](#) | [Portfolio](#)

EXPERIENCE

Siemens Healthineers

Software Engineer 1

Bangalore, India

July 2023 – Present

- **Built high-performance distributed microservices** in C++/Go supporting streaming data workloads; optimized concurrency and memory to reduce service latency by **25%**.
- **Implemented resilience patterns** like circuit breakers and retries, ensuring consistent system availability and data integrity for high-throughput compute infrastructure.
- **Enhanced multi-tenant data path** with MySQL tuning and Redis caching; increased ingestion throughput by **40%** for large-scale datasets.
- **Designed GPU-accelerated training pipelines** for ML/LLM workloads; implemented distributed strategies using PyTorch Distributed and DeepSpeed to scale across HPC clusters.
- **Optimized CUDA kernels** and compute flows for specialized hardware; performed deep performance profiling and memory analysis to eliminate bottlenecks in large-scale model training.
- **Supported containerized microservices** deployments on Linux/Azure and drove Agile delivery
- **Implemented idempotent ingestion handlers** with retries/backoff and deduplication keys to ensure reliable processing under failures; lowered duplicate processing by **90%** and improved recovery robustness.

Siemens Healthineers

Software Engineer Intern

Bangalore, India

February 2023 – July 2023

- Developed REST APIs and data models (C#) with MySQL/Redis; contributed to ingestion/validation modules and unit/integration tests.
- Automated developer workflows and CI tasks; participated in code/design reviews; achieved **60%** process time savings and **mentored** interns; adhered to Agile methodology.

SKILLS

Languages: C++, Python, CUDA, Golang, SQL, Bash, Shell, C#, Java

AI & Compute: PyTorch Distributed, DeepSpeed, HPC Clusters, Multi-GPU Systems, Parallel Processing, GenAI/LLM

Core Systems: Computer Architecture, Distributed Training Frameworks, Memory Management, RTOS, Linux Kernel

Optimization: CUDA Kernels, Performance Profiling (VTune, GDB), Lock-free Structures, High-Availability

Tools & Frameworks: PyTorch, TensorFlow, scikit-learn, Git, CI/CD, Docker, Kubernetes, AWS/Azure

PROJECTS

Distributed AI Compute Engine ([Repo](#)) | C++, CUDA, PyTorch

- Engineered a low-latency parallel compute engine in C++ using **lock-free data structures** and custom memory management to minimize overhead in distributed training tasks.
- Designed a **multithreaded** architecture for concurrent data processing; optimized for multi-node clusters and integrated resilient communication protocols for scalable model training.

Secure Distributed Training Utility ([Repo](#)) | C++, Python, DeepSpeed

- Built a secure utility for handling large-scale ML datasets with **AES-256** encryption; implemented block-level deduplication to optimize storage and data movement across HPC nodes.
- Crafted modular compute flows and robust error handling for data integrity; refined **CUDA kernel** performance for high-throughput training operations.

AI Compute Performance Analysis Tool ([Live](#) | [Repo](#)) | AI, ML, 3D Visualization

- Created an AI-powered platform that predicts compute risks, detects bottlenecks, and visualizes **GPU memory** dependencies; implemented ML-based performance hotspot prediction.
- Developed pipelines for repository analysis; contributed to **architecture/design reviews**; delivered dashboards to guide optimization of large-scale training jobs.

EDUCATION

National Institute of Engineering

Mysore, India

B.E in Electronics and Communication Engineering

August 2019 – July 2023

- Coursework: Computer Networks, Computer Architecture, DBMS, Big Data, Machine Learning, DSA, REST API.
- Language fluency: Kannada, English, Hindi.

CO-CURRICULARS

Xstasis Dance Group:Siemens Healthineers.| Project Head:UCSP Research Group, NIE.| Marketing Head:Onyx E-Cell, NIE.